

System Call Programming

Week 7

Processor Modes

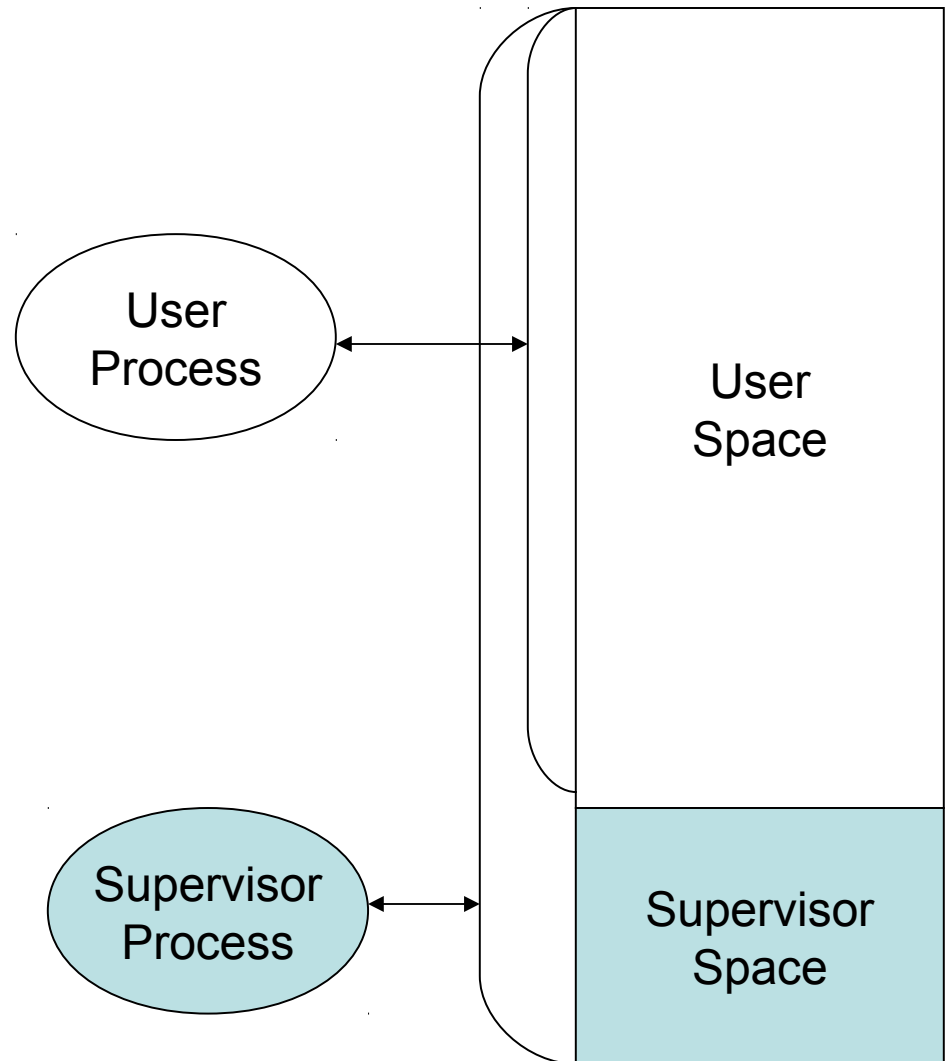
- Mode bit used to distinguish between execution on behalf of OS & execution on behalf of user.
- **Supervisor mode**: processor executes every instruction in its hardware repertoire.
- **User mode**: can only use subset of instructions

Processor Modes

- Instructions that can be executed in supervisor mode are supervisor privileges, or protection instruction
 - **I/O instructions** are protected. If an application needs to do I/O, it needs to get the OS to do it on its behalf.
 - Instructions that can change the **protection state** of the system are privileges (e.g., process' authorization status, pointers to resources, etc)

Processor Modes

- Mode bit may define areas of memory to be used when the processor is in supervisor mode vs user mode



The Kernel

- Core of OS software **executing in supervisor state**
- **Trusted software:**
 - Manages hardware resources (CPU, Memory and I/O)
 - Implements protection mechanisms that could not be changed through actions of untrusted software in user space
- **System call interface** is a **safe way** to expose privileged functionality and services of the processor

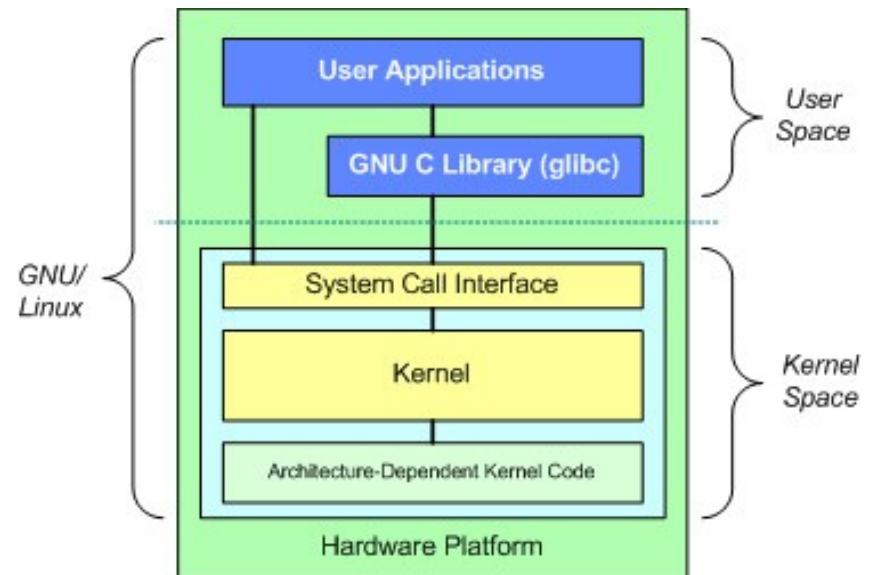
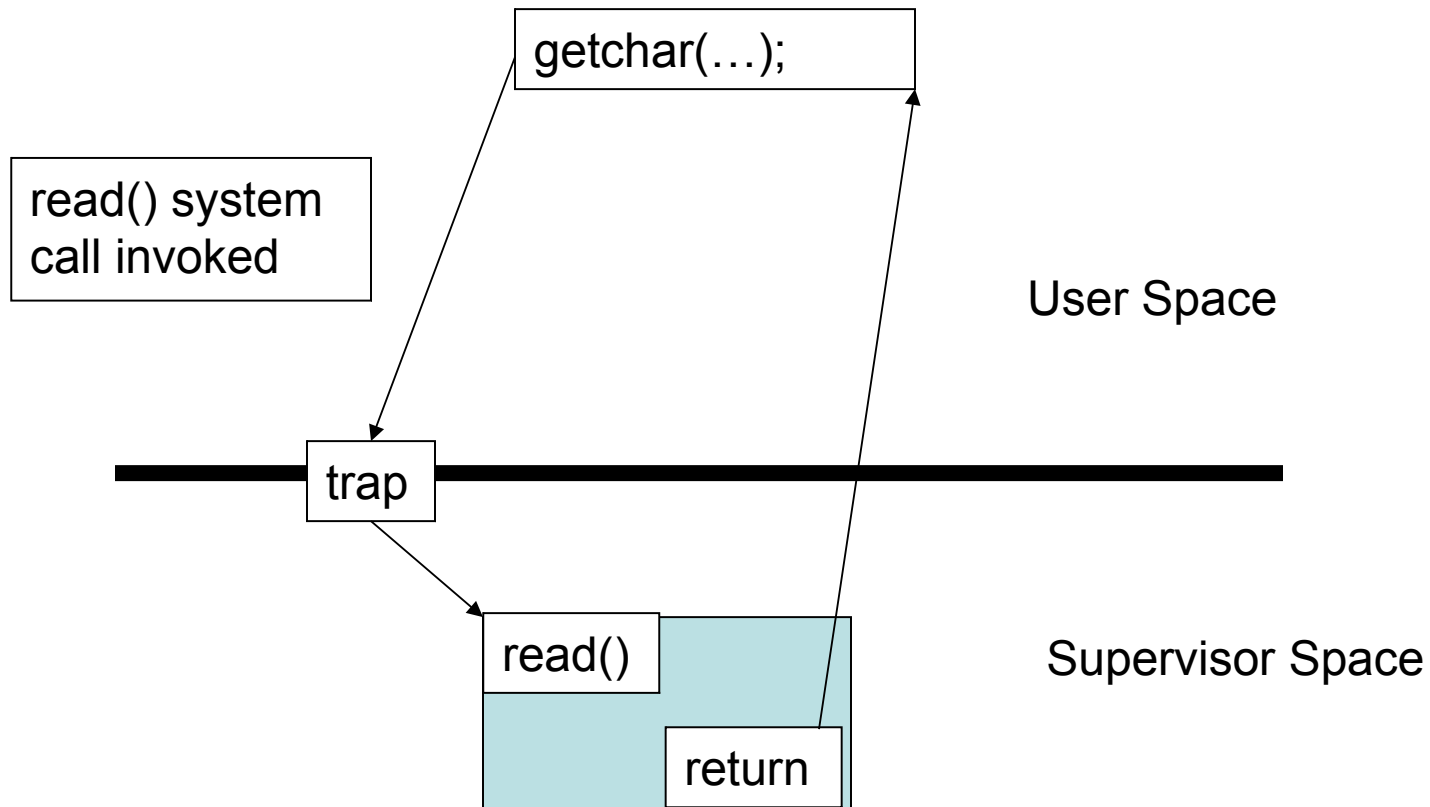


Image by: Tim Jones (IBM)

System Calls



Trap: System call causes a switch from user mode to kernel mode

System calls

- A system call involves the following
 - The system call causes a 'trap' that interrupts the execution of the user process (user mode)
 - The kernel takes control of the processor (kernel mode/privilege switch)
 - The kernel executes the system call on behalf of the user process
 - The user process gets back control of the processor (user mode/privilege switch)
- System calls have to be used **judiciously**. Expensive due to **privilege switching**.

System Calls

```
#include<unistd.h>
```

- `ssize_t read(int fildes, void *buf, size_t nbyte)`
 - *fildes*: file descriptor
 - *buf*: buffer to write to
 - *nbyte*: number of bytes to read
- `ssize_t write(int fildes, const void *buf, size_t nbyte);`
 - *fildes*: file descriptor
 - *buf*: buffer to write from
 - *nbyte*: number of bytes to write
- `int open(const char *pathname, int flags, mode_t mode);`
- `int close(int fd);`
- File descriptors
 - 0 stdin
 - 1 stdout
 - 2 stderr

Why are these system calls and not just regular library functions?

Example System Calls

- `pid_t getpid(void)`
 - Returns the process ID of the calling process
- `int dup(int fd)`
 - Duplicates a file descriptor `fd`. Returns a second file descriptor that points to the same file table entry as `fd` does.
- `int fstat(int filedes, struct stat *buf)`
 - Returns information about the file with the descriptor `filedes` into `buf`

```
struct stat {  
    dev_t      st_dev;      /* ID of device containing file */  
    ino_t      st_ino;      /* inode number */  
    mode_t     st_mode;     /* protection */  
    nlink_t    st_nlink;    /* number of hard links */  
    uid_t      st_uid;      /* user ID of owner */  
    gid_t      st_gid;      /* group ID of owner */  
    dev_t      st_rdev;     /* device ID (if special file) */  
    off_t      st_size;     /* total size, in bytes */  
    blksize_t  st_blksize;  /* blocksize for file system I/O */  
    blkcnt_t   st_blocks;   /* number of 512B blocks allocated */  
    time_t     st_atime;    /* time of last access */  
    time_t     st_mtime;    /* time of last modification */  
    time_t     st_ctime;    /* time of last status change */  
};
```

Laboratory

- Write a program using **getchar** and **putchar** to copy all bytes in stdin to stdout
- Write a program that uses **read** and **write** to read and write each byte, instead of using **getchar** and **putchar**. The **nbyte** argument should be 1 so it reads/writes a single byte at a time.
- Test it on a big file with 5000000 bytes

Laboratory

- read and write are system calls.
- getchar and putchar are library functions that live in user space but call read and write. And do buffering!
 - getchar() calls read() only if the local buffer is empty
 - getchar() reads multiple bytes into its buffer in one call to read()

Buffering Issues

- What is buffering?
- Why do we buffer?
- Can we make our buffer really big?

strace and time

- **strace:** Intercepts and prints out system calls to stderr or to an output file.

```
strace -o strace_output ./catb < bigfile.txt
```

- **time:** Timing

```
time ./catb < bigfile.txt
```

real	0m0.150s	(Total time elapsed)
user	0m0.132s	(time spent by CPU executing in user space)
sys	0m0.016s	(time spent by CPU executing in kernel space)

Buffered versus Unbuffered

- How many read and write calls are there in each?
- Why does one take longer than the other?

Writing to a File vs. Terminal

- Do you see any differences in the read/writes?
- In the buffered version, how many bytes are being read/written when outputting to a file and when outputting to the terminal?
- Why is there a difference?
- What are the performance considerations?

Homework

- Rewrite srot13 using system calls
- If stdin is a regular file, your program initially allocate enough memory to hold all data in the file all at once.
 - Otherwise, it should behave like before
- In addition to the functions you used last time, you'll need **read**, **write**, and **fstat**
- Read the man pages for these functions
- The program should be named **binsortu**
- Measure differences in performance between binsort and binsortu using the **time** command.