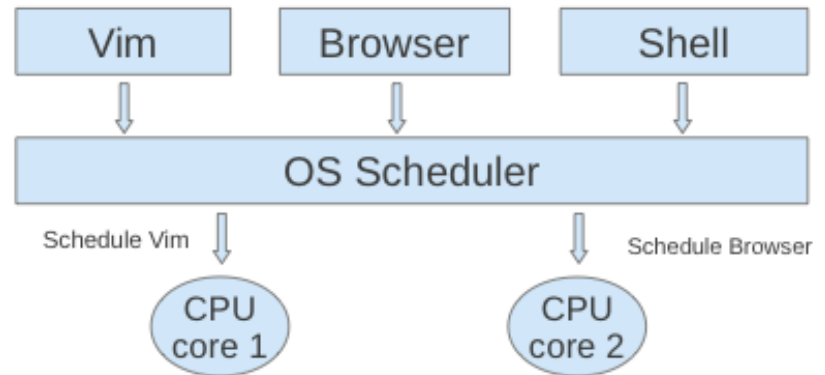
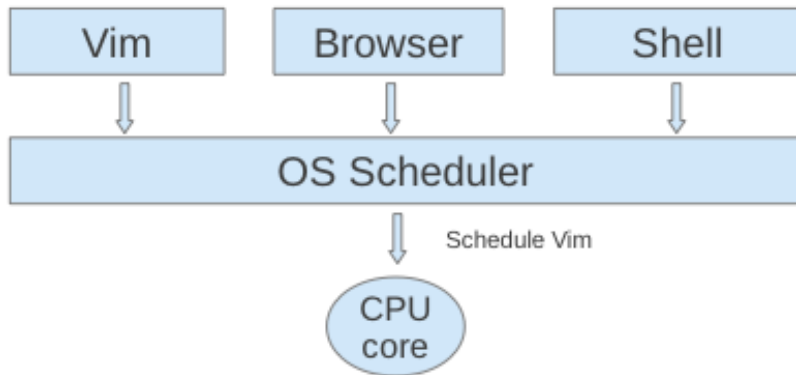


Multithreaded Performance

Week 9 - Monday

Multitasking

- Run multiple processes simultaneously to increase performance
- Processes do not share internal structures(stack, variable)
 - Communicate via IPC(inter-process communication) methods
 - Pipes,sockets, signals ,etc.
- Single core: Illusion of parallelism
- Multi-core: True parallelism



Multitasking

- `tr -s '[:space:]' '\n' | sort -u | comm -23 - words`
- Three separate processes spawned simultaneously
 - P1 `tr`
 - P2 `sort`
 - P3 `comm`
- Common buffers (pipes) exist between
- 2 processes for communication
 - ‘`tr`’ writes its `stdout` to a buffer that is read by ‘`sort`’
 - “`sort`” can execute when data is available in the buffer
 - Similarly, a buffer is used for communicating between ‘`sort`’ and ‘`comm`’

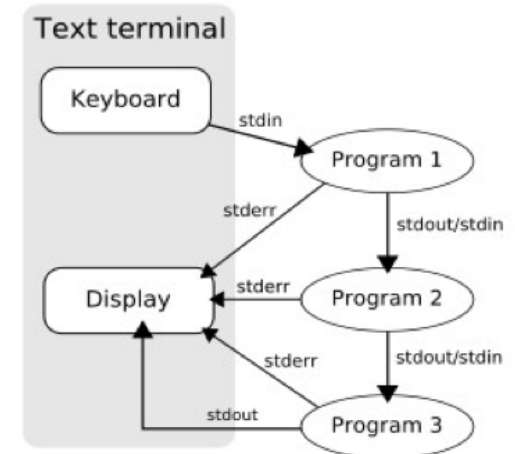


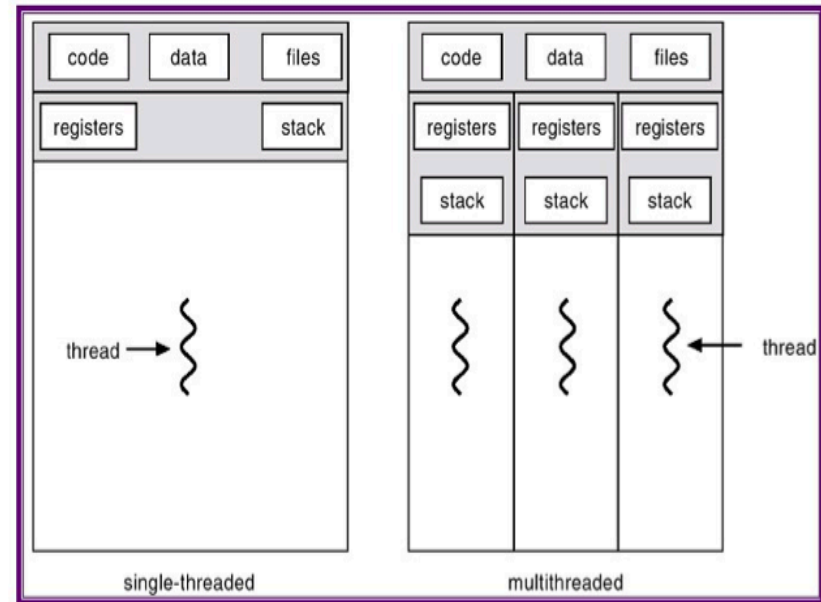
Image source: Wikipedia

Thread vs Process


- **Process:** An executing instance of an application
- **Thread:** Path of execution within a process
- A process can contain multiple threads
- When start an application, the OS creates a process and begins executing the primary thread of that process
- Note: Threads within the same process share the same address space, whereas different processes do not.
- This allows threads to read from and write to the same data structures and variables

Threads

- Memory
 - Private v/s Shared
- Each thread has its own
 - Stack
 - Registers
 - Thread ID
- Each thread shares the following with other threads in same process
 - Code
 - Global data
 - OS resources (files, I/O)
- Communication always through shared memory



Shared Memory Model



Threads

- Most common programming model for threads
 - Fork-join model

Intel® Software College

Relating Fork/Join to Code

The diagram illustrates the Fork/Join model. On the left, code snippets show sequential code and code with 'for' loops. In the center, a vertical line with two sets of three parallel vertical bars represents the fork and join points. On the right, labels indicate 'Sequential code' and 'Parallel code' sections.

Sequential code

Parallel code

Sequential code

Parallel code

Sequential code

55

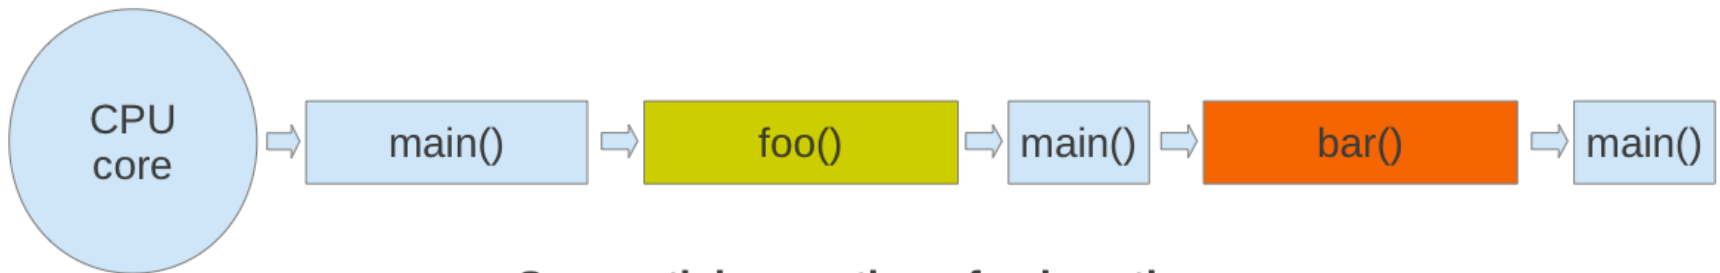
Shared-Memory Model and Threads

Copyright © 2006, Intel Corporation. All rights reserved.
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. *Other brands and names are the property of their respective owners.

Single threaded execution

```
int globalCounter = 0;  
  
int main() {  
  ...  
  foo(arg1, arg2);  
  bar(arg3, arg4, arg5);  
  ...  
  return 0;  
}
```

```
void bar(arg3, arg4, arg5) {  
  //code for bar  
}  
  
void foo(arg1, arg2) {  
  //code for foo  
}
```



Sequential execution of subroutines

Multi-threaded (single core)

```
int globalCounter = 0;

int main() {
...
Run thread foo(arg1, arg2);
Run thread bar(arg3, arg4, arg5);
...
return 0;
}
```

```
void bar(arg3, arg4, arg5) {
//code for bar
}

void foo(arg1, arg2) {
//code for foo
}
}
```



Time Sharing – Illusion of multithreaded parallelism
(Thread switching has less overhead compared to process switching)

Multi-threaded (multi core)


```
int globalCounter = 0;

int main() {
...
Run thread foo(arg1, arg2);
Run thread bar(arg3, arg4, arg5);
...
return 0;
}
```

```
void bar(arg3, arg4, arg5) {
//code for bar
}
```

```
void foo(arg1, arg2) {
//code for foo
}

}
```



True multithreaded parallelism

Multithreading properties

- Efficient way to parallelize tasks
- Thread switches are less expensive compared to process switches (context switching)
- Inter-thread communication is easy, via shared global data
- Need synchronization among threads accessing same data

PThread API

```
#include <pthread.h>
```

- `int pthread_create(pthread_t *thread, const pthread_attr_t *attr, void* (*thread_function)(void *), void *arg);`
 - Returns 0 on success, otherwise returns non-zero error number
- `void pthread_exit(void *retval);`
- `int pthread_join(pthread_t thread, void **retval);`
 - Returns 0 on success, otherwise returns non-zero error number

PThread example

```
#include<stdio.h>
#include<stdlib.h>
#include<pthread.h>
#include<unistd.h>

void* thread_fun(void *arg) {
    int thread_num = *(int *)arg;
    for (int i = 0; i < 5; i++) {
        printf("Thread #%-d: count %d\n",
            thread_num, i);
        sleep(1);
    }
    printf("Thread #%-d exit\n", thread_num);
    pthread_exit(NULL);
}
```

Gcc with `-lpthread` or `-pthread`!

```
int main() {
    pthread_t t1, t2;
    int *n1, *n2;
    n1 = malloc(sizeof(int));
    *n1 = 1;
    n2 = malloc(sizeof(int));
    *n2 = 2;
    printf("Create Thread #1\n");
    pthread_create(&t1, NULL,
        thread_fun, (void *)n1);
    printf("Create Thread #2\n");
    pthread_create(&t2, NULL,
        thread_fun, (void *)n2);
    printf("Finish creating threads\n");
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("Both threads have exited\n");
    return 0;
}
```

Thread Synchronization

- A simple example
 - A thread wait for another thread
 - `pthread_join()`
- The reality is much more complicated
 - For example, read/write shared data
 - Solution: all kinds of locks

Lab

- Environment variable for the right *sort*
 - Export PATH=/usr/local/cs/bin:\$PATH
- *od* to read random bytes
 - -A: how file offset are printed
 - -t: output format (translate bytes into different types)
 - -N: limit the # of bytes to read
- *tr/sed* to format the output
 - Replace space with newline
- Only *time* the *sort* program
 - Write the input of *sort* into a file
 - Write the output of *sort* into /dev/null

Homework

- Ray-Tracing: Powerful rendering technique in Computer Graphics
- Yields very high quality rendering
 - Suited for scenes with complex light interactions
 - Visually realistic
 - Trace the path of light in the scene
- Computationally very expensive
 - Not suited for rendering in real-time (example:games)
 - Suited for rendering high-quality pictures
- Embarrassingly parallel
 - Good candidate for multi-threading
 - Threads need not synchronize with each other, because each thread works on a different pixel